

Klion_Coding_HW3

March 29, 2022

```
[1]: %reload_ext autoreload
      %autoreload 2
```

```
[2]: import numpy as np
      import matplotlib.pyplot as plt
      import torch
      import starter
```

1 1. Feed Forward Neural network for Binary XOR

1.1 Part A:

```
[3]: def runAll(input_dim = 2, hidden_dim = 2, add_bias = False):
      model = starter.Model(input_dim, hidden_dim, add_bias = add_bias)
      B = starter.make_binary_arrays(input_dim)
      target = starter.generalize_xor(B)

      B = torch.from_numpy(B).float()
      target = torch.from_numpy(target).float()

      losses = starter.run_optimization(model, B, target, n_iter = 1000)
      return losses, model


      def sigmoid(z):
          return 1/(1 + np.exp(-z))

      def plot_(hidden, y_prob, points, weights, biases = (0, 0), segment = True):
          plt.scatter(hidden[:, 0], hidden[:, 1], c = y_prob)
          plt.scatter(points[:, 0], points[:, 1], c = ['white', 'black', 'black', 'white'], s = 50)
          if (segment):
              for weight, bias in zip(weights, biases):
                  line = starter.make_explicit_line_equation(weight, bias)
                  plt.plot(hidden[:, 0], line(hidden[:, 0]), color = 'black')
          return plt.get_current_fig_manager()
```

```
[4]: losses, model = runAll()
W1_np = model.W1.detach().numpy()
W2_np = model.W2.detach().numpy()
bias1_np = model.bias1.detach().numpy()
bias2_np = model.bias2.detach().numpy()
d_tensor = torch.Tensor([[0,0],
                          [0,1],
                          [1,0],
                          [1,1]])

plt.loglog(range(len(losses)), losses)
plt.ylabel("loss")
plt.xlabel("iteration")
plt.show()
```

```
100%|
| 1000/1000 [00:00<00:00, 2923.32it/s]
```



```
[5]: losses[-1]
```

```
[5]: 0.34871822595596313
```


Considering the loss does not go near 0, and gets stuck as iterations reach 1000, I would say the currently implemented trained model fails to implement this problem successfully.

2 Part B:

```
[6]: def makeGrid(x_int : tuple, y_int : tuple, num_points : int):  
    x = np.linspace(x_int[0], x_int[1], num = num_points)  
    y = np.linspace(y_int[0], y_int[1], num = num_points)  
    xv, yv = np.meshgrid(x, y)  
    merged = np.array((xv, yv))  
    flattened_tensor = torch.tensor(merged, dtype = torch.float).  
    ↪flatten(start_dim = 1).transpose(0, 1)  
    return flattened_tensor
```

```
[7]: x_interval, y_interval = (-0.5, 1.5), (-0.5, 1.5)  
grid_B = makeGrid(x_interval, y_interval, 100)  
model_y1, model_y2 = model(grid_B)  
model_y1 = model_y1.detach().numpy()  
model_y2 = model_y2.detach().numpy()
```


```
[8]: plt.rcParams['figure.figsize'] = (10,7)  
plot_(grid_B, model_y2, d_tensor, W1_np, biases = bias1_np, segment = False)  
plt.show()
```



$y = -x$ is one of the lines


2.1 Part C:

```
[9]: #plt.scatter(grid_B[:,0], grid_B[:, 1], c = model_y2)
#plt.scatter(0, 0, c = 0)
#plt.scatter(0, 1, c = 1)
#plt.scatter(1, 0, c = 1)
#plt.scatter(1, 1, c = 0)
#plt.plot(grid_B[:, 0], line(grid_B[:, 0]), color = 'black')
plot_(grid_B, model_y2, d_tensor, W2_np, biases = bias2_np, segment = True)
plt.xlim(-.1, 1.1)
plt.ylim(-.1, 1.1)
plt.show()
```




3 Part D:

```
[10]: y1_points, y2_points = model(d_tensor)
y1_points = y1_points.detach().numpy()
plot_(model_y1, model_y2, y1_points, W1_np, biases = bias1_np, segment = False)
plt.show()
```



```
[30]: plot_(sigmoid(model_y1), sigmoid(model_y2), sigmoid(y1_points), W1_np, biases = ↵
      ↪ bias1_np, segment = False)
      plt.show()
```




4 Part F:

```
[19]: losses_biased, model_biased = runAll(add_bias=True)
W1_biased_np = model_biased.W1.detach().numpy()
W2_biased_np = model_biased.W2.detach().numpy()
bias1_biased_np = model_biased.bias1.detach().numpy()
bias2_biased_np = model_biased.bias2.detach().numpy()

plt.loglog(range(len(losses_biased)), losses_biased)
plt.ylabel("loss")
plt.xlabel("iteration")
plt.show()
```

```
100%|
| 1000/1000 [00:00<00:00, 2600.95it/s]
```



```
[20]: losses_biased[-1]
```

```
[20]: 0.00044232996879145503
```

```
[23]: y1_biased, y2_biased = model_biased(grid_B)
y1_biased = y1_biased.detach().numpy()
y2_biased = y2_biased.detach().numpy()

y1_biased_points, y2_biased_points = model_biased(d_tensor)
y1_biased_points = y1_biased_points.detach().numpy()
```


```
[32]: plot_(grid_B, y2_biased, d_tensor, W1_biased_np, biases = bias1_biased_np,
↪segment = False)
plt.show()



plot_(grid_B, y2_biased, d_tensor, W1_biased_np, biases = bias1_biased_np)
plt.xlim(-.1, 1.1)
plt.ylim(-.1, 1.1)
plt.show()


plot_(y1_biased, y2_biased, y1_biased_points, W2_biased_np, biases =
↪bias2_biased_np, segment = False)
```

```
plt.show()

plot_(sigmoid(y1_biased), sigmoid(y2_biased), sigmoid(y1_biased_points), W2_biased_np, biases = bias2_biased_np)
plt.show()
```







5 Problem 2

6 Part A:

```
[16]: B3 = starter.make_binary_arrays(3)
      B3_xor = starter.generalize_xor(B3)
```

```
[39]: print(B3)
      print(B3_xor)
```

```
tensor([[0., 0., 0.],
        [0., 0., 1.],
        [0., 1., 0.],
        [0., 1., 1.],
        [1., 0., 0.],
        [1., 0., 1.],
        [1., 1., 0.],
        [1., 1., 1.]])
tensor([0., 1., 1., 0., 1., 0., 0., 1.]])
```

```
[33]: losses_3d, model_3d = runAll(input_dim = 3, add_bias = True)
```

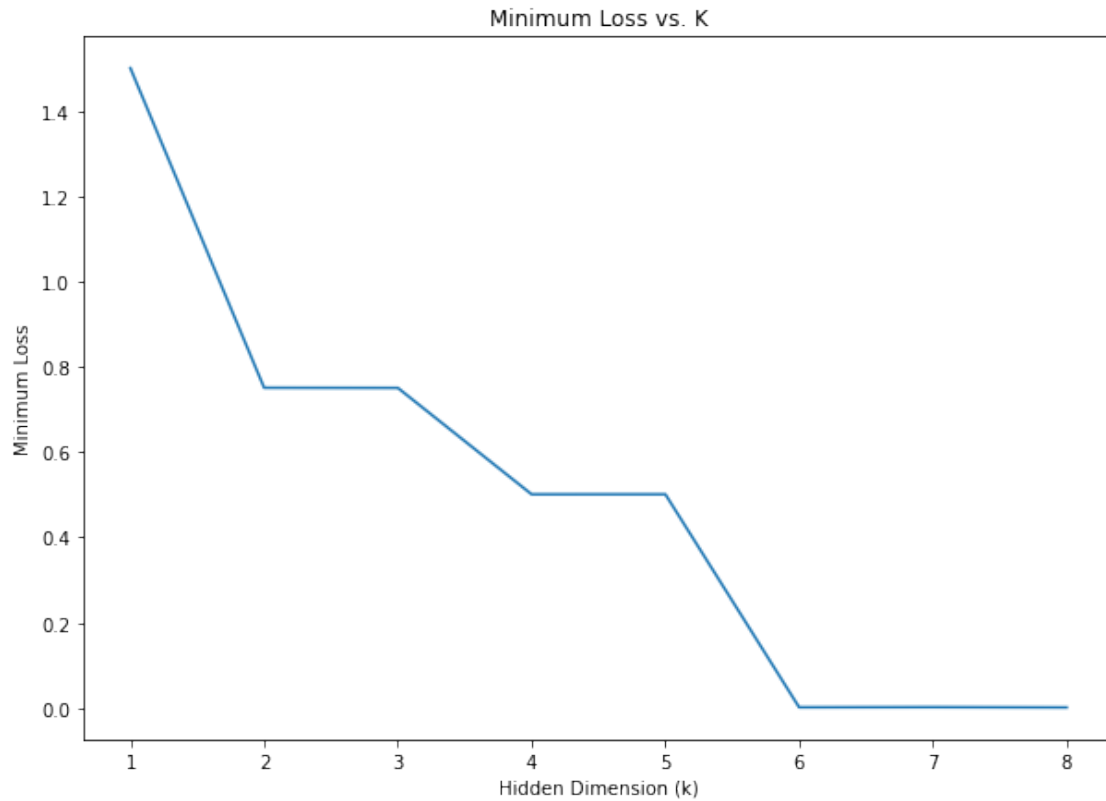
```
100%|  
| 1000/1000 [00:00<00:00, 3186.82it/s]
```

7 Part C:

```
[42]: ks = [1, 2, 3, 4, 5, 6, 7, 8]  
losses_k = []  
for k in ks:  
    local_loss, local_model = runAll(input_dim=3, hidden_dim = k, add_bias =  
    True)  
    losses_k.append(local_loss[-1])
```

```
100%|  
| 1000/1000 [00:00<00:00, 2750.69it/s]  
100%|  
| 1000/1000 [00:00<00:00, 2610.38it/s]  
100%|  
| 1000/1000 [00:00<00:00, 2596.81it/s]  
100%|  
| 1000/1000 [00:00<00:00, 2610.38it/s]  
100%|  
| 1000/1000 [00:00<00:00, 2593.38it/s]  
100%|  
| 1000/1000 [00:00<00:00, 2586.98it/s]  
100%|  
| 1000/1000 [00:00<00:00, 2570.42it/s]  
100%|  
| 1000/1000 [00:00<00:00, 2549.29it/s]
```

```
[45]: plt.plot(ks, losses_k)  
plt.xlabel("Hidden Dimension (k)")  
plt.ylabel("Minimum Loss")  
plt.title("Minimum Loss vs. K")  
plt.show()
```



It appears that the model starts to correctly classify the model at $k = 6$

8 Part E:

```
[51]: dims = [4, 5, 6]
      dims_losses = []
      for dim in dims:
          local_min_losses = []
          for k in ks:
              local_loss, local_model = runAll(input_dim = dim, hidden_dim = k,
          ↪add_bias = True)
              local_min_losses.append(local_loss[-1])
          print(local_min_losses)
          dims_losses.append(local_min_losses)
```


```
100%|
| 1000/1000 [00:00<00:00, 2796.55it/s]
100%|
| 1000/1000 [00:00<00:00, 2637.93it/s]
100%|
| 1000/1000 [00:00<00:00, 2603.58it/s]
```

100%|
 | 1000/1000 [00:00<00:00, 2630.98it/s]
 100%|
 | 1000/1000 [00:00<00:00, 2610.37it/s]
 100%|
 | 1000/1000 [00:00<00:00, 2627.50it/s]
 100%|
 | 1000/1000 [00:00<00:00, 2596.82it/s]
 100%|
 | 1000/1000 [00:00<00:00, 2603.58it/s]
 [4.0, 2.5461552143096924, 0.8053857684135437, 0.8158525228500366,
 0.8051002025604248, 0.01461966522037983, 0.037543900310993195,
 0.8037407398223877]
 100%|
 | 1000/1000 [00:00<00:00, 2769.47it/s]
 100%|
 | 1000/1000 [00:00<00:00, 2617.21it/s]
 100%|
 | 1000/1000 [00:00<00:00, 2179.15it/s]
 100%|
 | 1000/1000 [00:00<00:00, 2138.41it/s]
 100%|
 | 1000/1000 [00:00<00:00, 2282.59it/s]
 100%|
 | 1000/1000 [00:00<00:00, 2101.99it/s]
 100%|
 | 1000/1000 [00:00<00:00, 2118.16it/s]
 100%|
 | 1000/1000 [00:00<00:00, 2178.16it/s]
 [7.999999523162842, 8.0, 3.4433600902557373, 3.4526374340057373,
 3.4436800479888916, 3.4497451782226562, 3.448378562927246, 0.8758033514022827]
 100%|
 | 1000/1000 [00:00<00:00, 2590.09it/s]
 100%|
 | 1000/1000 [00:00<00:00, 2216.80it/s]
 100%|
 | 1000/1000 [00:00<00:00, 2236.63it/s]
 100%|
 | 1000/1000 [00:00<00:00, 2251.74it/s]
 100%|
 | 1000/1000 [00:00<00:00, 2093.67it/s]
 100%|
 | 1000/1000 [00:00<00:00, 2118.17it/s]
 100%|
 | 1000/1000 [00:00<00:00, 2178.16it/s]
 100%|

| 1000/1000 [00:00<00:00, 2048.72it/s]

[16.0, 10.785126686096191, 9.929214477539062, 9.922440528869629,
9.945112228393555, 4.444655895233154, 9.924810409545898, 4.465934753417969]

```
[58]: plt.plot(ks, dims_losses[0], label = 'd = 4')  
plt.plot(ks, dims_losses[1], label = 'd = 5')  
plt.plot(ks, dims_losses[2], label = 'd = 6')  
plt.xlabel("Hidden Dimension (k)")  
plt.ylabel("Minimum Loss")  
plt.title("Minimum Loss vs. K")  
plt.legend()  
plt.show()
```



6 hyperplanes appear to successfully classify the d-dimensional problem