

Homework 8, due March 2nd, 11:59pm

February 23, 2022

1. Download the files `cnnttrain.zip` and `cnntest.zip` from Canvas, containing training and test patches of size 64×64 for the problem of predicting the resolution of an image. The first two digits in each file name represent the resolution y_i of the patch (between 10% and 96%).

The goal of this project is to train a regression CNN to predict the resolution. We will use the square loss functions on the training examples $(\mathbf{x}_i, y_i), i = 1, \dots, n$:

$$S(\mathbf{w}) = \frac{1}{n} \sum_{i=1}^n (y_i - f_{\mathbf{w}}(\mathbf{x}_i))^2 + \lambda \|\mathbf{w}\|^2 \quad (1)$$

Besides the loss function, we will measure the R^2 , defined as:

$$R^2 = 1 - \frac{\sum_{i=1}^n (y_i - \hat{y}_i)^2}{\sum_{i=1}^n (y_i - \bar{y})^2}$$

where $\hat{y}_i = f_{\mathbf{w}}(\mathbf{x}_i)$ and $\bar{y} = \frac{1}{n} \sum_{i=1}^n y_i$.

Experiment with different CNN architectures to obtain a good result. One example of a CNN you could use contains five convolutional layers with stride 1 and no padding, the first four with filters of size 5×5 with or without holes (dilation), and the last layer of the appropriate size to obtain a 1×1 output. The first two convolutions have 16 filters, the next two have 32 filters, and the last has only one filter. The second and third convolutions could be followed by 2×2 max pooling with stride 2 respectively. The fourth convolution layer is followed by ReLU.

Try to use a GPU and CUDA for faster training.

- Train a CNN for 100 epochs with momentum 0.9 using the square loss (1). Use the SGD optimizer with an appropriate learning rate and $\lambda = 0.0001$ (weight decay). Start with minibatch size 64 and double it every 20 epochs and try to obtain a good training R^2 (at least 0.9). You might also need to reduce the learning rate as training progresses. Show a plot of the loss function vs epoch number for the training set and the test set. Show another plot of the training and test R^2 vs epoch number. (4 points)
- Repeat point a) with the same CNN architecture, but with a fixed minibatch of 1024 and a fixed learning rate. It's ok if you cannot get a training R^2 of at least 0.9 in this case. (3 points)
- Report the CNN architecture in a table, where each row describes one layer, including the layer description, size and number of parameters, and the last row containing the total number of parameters. (2 points)
- Plot with two different colors the train and test residuals $r_i = y_i - \hat{y}_i$ vs y_i for the model obtained at a). (1 point)