

March 23rd, 2022

1.

a. Most likely sequence y

[illegible]

b. $\alpha_{128}^1 / \alpha_{128}^2 = 1.634322846$

c. $\beta_{128}^1 / \beta_{128}^2 = 0.800134907$

2. $\pi = [0.66354372 \quad 0.33645628]$
$$a = \begin{bmatrix} 0.53691227 & 0.46308773 \\ 0.47678048 & 0.52321592 \end{bmatrix}$$
$$b = \begin{bmatrix} 0.1826396 & 0.2165343 & 0.1756936 & 0.2094443 & 0.1155081 & 0.1001801 \\ 0.2110478 & 0.1872504 & 0.2015701 & 0.1811477 & 0.1328683 & 0.0861157 \end{bmatrix}$$

```

#!/usr/bin/env python
# coding: utf-8

# In[780]:

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

# # Viterbi

# In[782]:

# From "https://stackoverflow.com/questions/9729968/python-implementation-of-viterbi-
algorithm"
def viterbi(y, A, B, Pi=None):
    """
    Return the MAP estimate of state trajectory of Hidden Markov Model.

    Parameters
    -----
    y : array (T,)
        Observation state sequence. int dtype.

    A : array (K, K)
        State transition matrix. See HiddenMarkovModel.state_transition for
        details.

    B : array (K, M)
        Emission matrix. See HiddenMarkovModel.emission for details.

    Pi: optional, (K,)
        Initial state probabilities: Pi[i] is the probability  $x[0] == i$ . If
        None, uniform initial distribution is assumed ( $Pi[:] == 1/K$ ).

```

Returns

x : array (T,)

Maximum a posteriori probability estimate of hidden state trajectory,
conditioned on observation sequence y under the model parameters A, B,
Pi.

T1: array (K, T)

the probability of the most likely path so far

T2: array (K, T)

the x_{j-1} of the most likely path so far

"""

Cardinality of the state space

K = A.shape[0]

Initialize the priors with default (uniform dist) if not given by caller

Pi = Pi if Pi is not None else np.full(K, 1 / K)

T = len(y)

T1 = np.empty((K, T), 'd')

T2 = np.empty((K, T), 'B')

Initialize the tracking tables from first observation

T1[:, 0] = np.log(Pi * B[:, y[0]])

T2[:, 0] = 0

Iterate through the observations updating the tracking tables

for i in range(1, T):

#need to subtract 1 from y as it ranges from 1-6 but indexed 0 in array

T1[:, i] = np.max(T1[:, i - 1] + np.log(A), 1) + np.log(B[np.newaxis, :, y[i] -
1])

T2[:, i] = np.argmax(T1[:, i - 1] + np.log(A), 1)

Build the output, optimal model trajectory

x = np.empty(T, 'B')

```

x[-1] = np.argmax(T1[:, T - 1])
#flip since we are backtracking
for i in reversed(range(1, T)):
    x[i - 1] = T2[x[i], i]

return x, T1, T2

```

```
# In[783]:
```

```

obs_a = np.loadtxt("../datasets/markov/hmm_pb1.csv", dtype = int, delimiter = ",")
obs_a_pd = pd.read_csv("../datasets/markov/hmm_pb1.csv", header=None)
pi_a = np.array([0.5, 0.5])
trans_a = np.array([
    [0.95, 0.05],
    [0.05, 0.95]
])
emit_a = np.array([
    [1/6, 1/6, 1/6, 1/6, 1/6, 1/6],
    [1/10, 1/10, 1/10, 1/10, 1/10, 1/2]
])

```

```
# In[1186]:
```

```
y, T1, T2 = viterbi(obs_a, trans_a, emit_a, pi_a)
```

```
# In[1188]:
```

```
#y = 1 is Fair, y=2 is loaded
```

```
print(y+1)
print(T1)
print(T2)
```

```
# # Forward Algorithm
```

```
# In[1216]:
```

```
#Modified from "http://www.adeveloperdiary.com/data-science/machine-learning/forward-
and-backward-algorithm-in-hidden-markov-model/"
```

```
def forward(observations, transition_matrix, emission_matrix, pi):
    #python indexed 0 so subtract 1 from all dice rolls
    observations = observations - 1
    #get number of observations
    observations_rows = observations.shape[0]
    #get number of states (K)
    K = transition_matrix.shape[0]

    #Initialize alpha probabilities
    alpha = np.zeros([observations_rows, K])
    alpha[0, :] = pi * emission_matrix[:, observations[0]]
    alpha_probs = np.zeros([observations_rows, K])
    alpha_probs[0] = alpha[0] / np.sum(alpha[0])

    #common factor to avoid underflow
    #u = np.zeros([observations_rows - 1, K])

    #iterate forward through observations to compute alpha probabilities
    for t, obs in enumerate(observations[1:], 1):
        alpha[t] = np.array([
            #start with probability we go from previous state to fair
            emission_matrix[0][obs] * np.sum([
```

```

        #chance from fair to fair
        alpha[t - 1][0] * transition_matrix[0, 0],
        #chance from loaded to fair
        alpha[t - 1][1] * transition_matrix[1, 0]
    ]),
    #go from previous state to loaded
    emission_matrix[1][obs] * np.sum([
        #chance from fair to loaded
        alpha[t - 1][0] * transition_matrix[0, 1],
        #chance from loaded to loaded
        alpha[t - 1][1] * transition_matrix[1, 1]
    ]),
    ])
alpha_probs[t] = np.array([
    #start with probability we go from previous state to fair
    emission_matrix[0][obs] * np.sum([
        #chance from fair to fair
        alpha_probs[t - 1][0] * transition_matrix[0, 0],
        #chance from loaded to fair
        alpha_probs[t - 1][1] * transition_matrix[1, 0]
    ]),
    #go from previous state to loaded
    emission_matrix[1][obs] * np.sum([
        #chance from fair to loaded
        alpha_probs[t - 1][0] * transition_matrix[0, 1],
        #chance from loaded to loaded
        alpha_probs[t - 1][1] * transition_matrix[1, 1]
    ]),
    ])
#normalize for probabilities
alpha_probs[t] /= np.sum(alpha_probs[t])
#still need u_t for underflow ?
#u[t - 1, :] = alpha[t] / alpha[t - 1]

```

```
print(alpha[128][0] / alpha[128][1])
return alpha_probs
```

```
# In[1214]:
```

```
alpha = forward(obs_a, trans_a, emit_a, pi_a)
```

```
# In[1215]:
```

```
print(alpha)
```

```
# # Backwards Algorithm
```

```
# In[1217]:
```

```
#Modified from from "http://www.adeveloperdiary.com/data-science/machine-
learning/forward-and-backward-algorithm-in-hidden-markov-model/"
```

```
def backward(observations, transition_matrix, emission_matrix):
    #python indexed 0 so subtract 1 from all dice rolls
    observations = observations - 1
    #get number of observations
    observations_rows = observations.shape[0]
    #get number of states
    K = transition_matrix.shape[0]

    #initial probabilities for beta start at T and work backwards
    beta, beta_probs = np.zeros([observations_rows, K]), np.zeros([observations_rows,
K])
```

```

beta[observations_rows - 1] = np.array([1, 1])
beta_probs[observations_rows - 1] = np.array([1, 1])

for t, obs in reversed(list(enumerate(observations))[1:]):
    beta[t - 1] = np.array([
        np.sum([
            #from fair to fair
            transition_matrix[0, 0] * emission_matrix[0][obs] * beta[t][0],
            #from fair to loaded
            transition_matrix[0, 1] * emission_matrix[1][obs] * beta[t][1],
        ]),
        np.sum([
            #from loaded to fair
            transition_matrix[1, 0] * emission_matrix[0][obs] * beta[t][0],
            #from loaded to loaded
            transition_matrix[1, 1] * emission_matrix[1][obs] * beta[t][1]
        ])
    ])
    beta_probs[t - 1] = np.array([
        np.sum([
            #from fair to fair
            transition_matrix[0, 0] * emission_matrix[0][obs] * beta_probs[t][0],
            #from fair to loaded
            transition_matrix[0, 1] * emission_matrix[1][obs] * beta_probs[t][1],
        ]),
        np.sum([
            #from loaded to fair
            transition_matrix[1, 0] * emission_matrix[0][obs] * beta_probs[t][0],
            #from loaded to loaded
            transition_matrix[1, 1] * emission_matrix[1][obs] * beta_probs[t][1]
        ])
    ])

#normalize beta values

```



```
        beta_probs[t - 1] /= np.sum(beta_probs[t - 1])
    print(beta[128][0] / beta[128][1])

    return beta_probs
```

```
# In[1218]:
```

```
beta = backward(obs_a, trans_a, emit_a)
```

```
# In[1219]:
```

```
print(beta)
```

```
# # 2. Baum Welch Algorithm
```

```
# In[958]:
```

```
obs_b = pd.read_csv("../datasets/markov/hmm_pb2.csv", header=None).values
obs_b = obs_b[0]
trans_EM = np.random.rand(2, 2)
trans_EM /= trans_EM.sum(axis=1)[:, None]
emit_EM = np.random.rand(2, 6)
emit_EM /= emit_EM.sum(axis=1)[:, None]
pi_EM = np.random.rand(2)
pi_EM /= pi_EM.sum()

#Check initial guesses
print(trans_EM)
```

```

print(emit_EM)
print(pi_EM)
#Check their probabilities add to 1
print(trans_EM.sum(axis=1))
print(emit_EM.sum(axis=1))
print(np.sum(pi_EM))

# In[1221]:

def baum_welch(observations, transition_matrix, emission_matrix, pi, epochs = 50):

    observations = observations - 1
    observations_rows = observations.shape[0]
    #number of states
    K = transition_matrix.shape[0]

    for epoch in range(epochs):
        #initialize placeholders for observations
        new_transition_matrix = np.zeros((2, 2))
        new_pi = np.zeros(2)
        new_emission_matrix = np.zeros((2, 6))
        fitness = 0

        alphas = forward(observations, transition_matrix, emission_matrix, pi)
        betas = backward(observations, transition_matrix, emission_matrix)

        #collect probabilities of each observation
        probability_of_obs = pi[0] * betas[0][0] + pi[1] * betas[0][1]
        fitness += probability_of_obs

        #xi calculations - probability of being in state 'i' at time 't'
        #and state 'j' at time 't+1' based on the model

```

```

xi = np.zeros([observations_rows, K, K])
for t in range(observations_rows - 1):
    denominator = np.dot(np.dot(alphas[t, :].T, transition_matrix) *
emission_matrix[:, observations[t + 1]].T, betas[t + 1, :])
    xi[t, 0, 0] = (alphas[t][0] * transition_matrix[0, 0]
                    * emission_matrix[0][observations[t + 1]] * betas[t + 1][0]) /
denominator
    xi[t, 0, 1] = (alphas[t][0] * transition_matrix[0, 1]
                    * emission_matrix[1][observations[t + 1]] * betas[t + 1][1]) /
denominator
    xi[t, 1, 0] = (alphas[t][1] * transition_matrix[1, 0]
                    * emission_matrix[0][observations[t + 1]] * betas[t + 1][0]) /
denominator
    xi[t, 1, 1] = (alphas[t][1] * transition_matrix[1, 1]
                    * emission_matrix[1][observations[t + 1]] * betas[t + 1][1]) /
denominator

#gamma stores probability that we are in state 'i' at time t
gamma = np.zeros([observations_rows, K])
for t in range(observations_rows - 1):
    gamma[t, 0] = np.sum(xi[t, 0, :])
    gamma[t, 1] = np.sum(xi[t, 1, :])

#recompute initial probabilities for the observation
new_pi = gamma[0] / np.sum(gamma[0])

# calculate new transition matrix for later update
new_transition_matrix[0, 0] += np.sum(xi[:, 0, 0]) / np.sum(gamma[:, 0])
new_transition_matrix[0, 1] += np.sum(xi[:, 0, 1]) / np.sum(gamma[:, 0])
new_transition_matrix[1, 0] += np.sum(xi[:, 1, 0]) / np.sum(gamma[:, 1])
new_transition_matrix[1, 1] += np.sum(xi[:, 1, 1]) / np.sum(gamma[:, 1])

# calculate new emission matrix for later update
emit_denom = np.sum(gamma, axis=0)
for l in range(emission_matrix.shape[1]):

```

```
        #new_emission_matrix[:, l] = np.sum(gamma[observations == l, :]) /  
emit_denom
```

```
        new_emission_matrix[0, l] = np.sum(np.take(gamma[:, 0],  
np.where(observations == l))) / emit_denom[0]
```

```
        new_emission_matrix[1, l] = np.sum(np.take(gamma[:, 1],  
np.where(observations == l))) / emit_denom[1]
```

```
    #update all values for probabilities for next epoch
```

```
    pi = new_pi
```

```
    transition_matrix = new_transition_matrix
```

```
    emission_matrix = new_emission_matrix
```

```
# Some logging is always good :)
```

```
    print('EPOCH #{} ='.format(epoch))
```

```
    print('Transition Matrix:', transition_matrix)
```

```
    print('Initial Dice Probability:', pi)
```

```
    print('Emission Matrix:', emission_matrix)
```

```
    print('Fitness:', fitness)
```

```
    print()
```

```
# In[1183]:
```

```
baum_welch(obs_b, trans_EM, emit_EM, pi_EM, epochs = 1000)
```